MICRO PROJECT REPORT

OLAP Operations: Load Business Data and Create OLAP Cubes using Apache Superset

Submitted by:
Chetan D R — 2023BIFT07AED039
Ashu Singh — 2023BIFT07AED007
Arpita Sinha — 2023BIFT07AED022

In partial fulfilment for the award of the degree of

BACHELOR OF TECHNOLOGY
In
INFORMATION TECHNOLOGY

Under the Guidance of
Dr. Vinita Nair
Associate Professor, ASAC

Alliance School of Advanced Computing
Alliance University, Bengaluru
October 2025

INFORMATION TECHNOLOGY
ALLIANCE SCHOOL OF ADVANCED COMPUTING

CERTIFICATE

This is to certify that the report on Information Technology in Integration Systems is submitted by Chetan D R (Reg. No. 2023BIFT07AED039), Ashu Singh (Reg. No. 2023BIFT07AED007), and Arpita Sinha (Reg. No. 2023BIFT07AED022) in partial fulfillment for the award of the degree of Bachelor of Technology (Information Technology) of Alliance University.

This is a bonafide work accomplished under my supervision and guidance during the academic year 2026. This report embodies the results of original work and studies conducted by the students, and the contents do not form the basis for the award of any other degree to the candidate or to any other person.

Dr. Vinita Nair
Associate Professor, Alliance School of Advanced Computing

INFORMATION TECHNOLOGY
ALLIANCE SCHOOL OF ADVANCED COMPUTING

DECLARATION

I hereby declare that the project report titled "Business Data Analysis using OLAP and Apache Superset" is a bonafide work carried out by Chetan D R (Reg. No. 2023BIFT07AED039), Ashu Singh (Reg. No. 2023BIFT07AED007), and Arpita Sinha (Reg. No. 2023BIFT07AED022) in partial fulfillment for the award of the degree of Bachelor of Technology (Information Technology) of Alliance University during the academic year 2026.

This report has not been submitted for the award of any other degree or diploma in any institution. The work presented herein is original and has been carried out independently by us to the best of our knowledge.

Name of the Student

Chetan D R — 2023BIFT07AED039
Ashu Singh — 2023BIFT07AED007
Arpita Sinha — 2023BIFT07AED022

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to everyone who contributed to the successful completion of this project titled "Business Data Analysis using OLAP and Apache Superset." This project has provided valuable knowledge and practical experience in the field of data analysis, business intelligence, and modern data visualization tools, which will be beneficial throughout our careers.

I would like to extend my heartfelt thanks to our mentor and advisor, Prof. Vinita Nair, for her constant guidance, encouragement, and valuable feedback during the development of this project. Her support and expertise played a crucial role in the successful completion of this work.

I am also thankful to the professors of the Alliance School of Advanced Computing for providing the necessary knowledge, resources, and a supportive academic environment that enabled us to carry out this project effectively. Their teaching has helped us gain a deeper understanding of advanced database technologies and data analytics.

Through this project, we gained hands-on experience in concepts such as OLAP operations, ETL processes, and data visualization using Apache Superset. This experience enhanced our understanding of how businesses analyze large datasets and make data-driven decisions.

Thank you.

ABSTRACT

In today's data-driven world, organizations generate large volumes of data that require efficient analysis to support decision-making. Traditional database systems are not always sufficient for performing complex analytical queries. This project focuses on Business Data Analysis using OLAP (Online Analytical Processing) and Apache Superset, providing an efficient way to analyze and visualize large datasets.

The system is designed to use the ETL (Extract, Transform, Load) process to collect data from multiple sources, clean and transform it into a consistent format, and store it in a centralized system for analysis. Using OLAP techniques, data is organized into multidimensional structures known as cubes, enabling operations such as slice, dice, roll-up, drill-down, and pivot for in-depth analysis.

To enhance data visualization and user interaction, Apache Superset is used to create dynamic dashboards and charts. These dashboards help users identify patterns, trends, and key insights from business data, such as sales performance, regional distribution, and product analysis.

The project demonstrates how modern business intelligence tools and data warehousing techniques can be integrated to provide meaningful insights and improve decision-making. Overall, this system highlights the importance of combining OLAP operations with powerful visualization tools to make data analysis more efficient, interactive, and user-friendly.

TABLE OF CONTENTS

1. Introduction 1
2. Problem Statement 2
3. Objectives 3
4. Methodology 4
5. Hardware Requirements 5
6. Software Requirements 5
7. System Architecture 6
8. Implementation 7
9. Results and Output 8
   9.1 Dataset View 8
   9.2 Roll-Up Operation 9
   9.3 Drill-Down Operation 10
   9.4 Slice Operation 11
   9.5 Dice Operation 12
   9.6 Bar Graph Visualization 13
   9.7 Pie Chart Visualization 14
   9.8 Pivot Operation 15
10. Conclusion 16
11. Future Scope 16
12. References 17

1. INTRODUCTION

In today's digital era, organizations generate massive amounts of data from various sources such as transactions, customer interactions, and business operations. Analyzing this data effectively is essential for making informed business decisions. However, traditional database systems are not optimized for complex analytical queries and large-scale data analysis.

To address this challenge, Online Analytical Processing (OLAP) has emerged as an important technology that enables fast and efficient analysis of multidimensional data. OLAP allows users to view data from different perspectives using operations such as slice, dice, roll-up, drill-down, and pivot, making it easier to identify patterns, trends, and insights.

This project, titled "Business Data Analysis using OLAP and Apache Superset," focuses on implementing OLAP concepts along with modern data visualization tools to analyze business data effectively. The system uses the ETL (Extract, Transform, Load) process to collect data from multiple sources, transform it into a consistent format, and store it in a centralized database for analysis.

To enhance data interpretation and visualization, Apache Superset is used as a business intelligence tool. It enables the creation of interactive dashboards and charts such as bar graphs, line charts, and pie charts, helping users understand complex data in a simple and visual manner.

The project also demonstrates how businesses can analyze sales data based on different dimensions such as time, location, and product. This helps in identifying trends, improving decision-making, and optimizing business performance.

Overall, this project highlights the importance of combining OLAP techniques with powerful visualization tools to transform raw data into meaningful insights and support data-driven decision-making in modern organizations.

2. PROBLEM STATEMENT

In today's data-driven environment, organizations generate large volumes of data from multiple sources such as sales, customer transactions, and operational activities. However, traditional database systems are not designed to efficiently handle complex analytical queries and multidimensional data analysis.

Most existing systems focus on storing and retrieving data rather than providing meaningful insights. As a result, businesses face difficulties in analyzing data across different dimensions such as time, location, and products. This limits their ability to identify trends, patterns, and key performance indicators.

Additionally, raw data is often scattered across multiple sources and formats, making it inconsistent and difficult to analyze. Without proper data integration and transformation processes, decision-making becomes slow and less effective.

Furthermore, many organizations lack user-friendly visualization tools that can present data in an interactive and understandable format. This makes it challenging for non-technical users to interpret data and derive actionable insights.

Therefore, there is a need for a system that can efficiently integrate data from multiple sources, perform multidimensional analysis using OLAP techniques, and provide interactive visualizations through tools like Apache Superset to support better and faster decision-making.

3. OBJECTIVES

The primary objectives of this project are as follows:

- To understand and implement OLAP (Online Analytical Processing) concepts for efficient multidimensional data analysis.
- To design and implement an ETL (Extract, Transform, Load) process for integrating data from multiple sources into a centralized system.
- To perform data analysis using OLAP operations such as slice, dice, roll-up, drill-down, and pivot.
- To develop interactive and user-friendly data visualizations using Apache Superset.
- To analyze business data (such as sales, products, and regions) in order to identify trends and patterns.

4. METHODOLOGY

The development of the project "Business Data Analysis using OLAP and Apache Superset" follows a structured approach to ensure efficient data integration, analysis, and visualization.

Initially, relevant business data (such as sales, product, and regional data) is collected from different sources. This data may exist in different formats and structures, making it unsuitable for direct analysis.

To address this, the ETL (Extract, Transform, Load) process is implemented. In the extraction phase, data is gathered from multiple sources. In the transformation phase, the data is cleaned, formatted, and converted into a consistent structure. Finally, in the loading phase, the processed data is stored in a centralized database or data warehouse.

Once the data is prepared, OLAP (Online Analytical Processing) techniques are applied. The data is organized into multidimensional structures known as cubes, consisting of dimensions (such as time, location, and product) and measures (such as sales and revenue).

Various OLAP operations are then performed to analyze the data effectively:

- Slice — to focus on a single dimension
- Dice — to filter multiple dimensions
- Roll-up — to summarize data at a higher level
- Drill-down — to view detailed data
- Pivot — to change the perspective of analysis

For data visualization, Apache Superset is used. The processed data is connected to Superset, where datasets are created and used to generate charts such as bar graphs, line charts, and pie charts.

Overall, this methodology ensures a systematic flow from raw data collection to meaningful data visualization and analysis.

5. HARDWARE REQUIREMENTS

The OLAP-based system is primarily software-dependent and does not require specialized hardware. The following standard hardware configuration is sufficient to run the system effectively:

- Computer or Laptop
- Processor: Intel Core i3 or higher
- RAM: Minimum 4 GB (8 GB recommended for large datasets)
- Storage: At least 10 GB of free disk space

6. SOFTWARE REQUIREMENTS

The project uses the following software tools and technologies:

- Python / SQL — Used for data processing, querying, and performing analytical operations on the dataset.
- MySQL / PostgreSQL — Used as the database management system to store structured business data such as sales, products, and regional information.
- Apache Superset — Used as the data visualization and business intelligence tool to create dashboards, charts, and reports.
- ETL Tools or Scripts — Used to perform Extract, Transform, and Load operations for integrating and preparing data from multiple sources.
- Web Browser (Chrome / Edge) — Used to access Apache Superset dashboards and visualize data.
- Visual Studio Code / Jupyter Notebook — Used as the development environment for writing queries, scripts, and managing the project.

7. SYSTEM ARCHITECTURE

The system architecture of the project "Business Data Analysis using OLAP and Apache Superset" consists of three main layers that work together to perform data integration, analysis, and visualization.

7.1 Data Source Layer

This layer includes multiple data sources such as business datasets covering sales data, product details, and regional data. The data may originate from different systems and exist in various formats such as CSV files or relational databases. These raw data sources form the foundation for all subsequent analysis.

7.2 Data Processing Layer (ETL + OLAP)

In this layer, the ETL (Extract, Transform, Load) process is applied:

- Extract — Data is collected from multiple sources.
- Transform — Data is cleaned, formatted, and converted into a consistent structure.
- Load — Processed data is stored in a centralized database or data warehouse.

After ETL, OLAP (Online Analytical Processing) is applied to organize data into multidimensional cubes. These cubes enable efficient analysis through operations such as Slice, Dice, Roll-up, Drill-down, and Pivot. This layer is responsible for converting raw data into structured and meaningful information.

7.3 Visualization Layer (Apache Superset)

This layer uses Apache Superset as a business intelligence tool to visualize the processed data. The system connects Superset to the database, where datasets are created and used to build charts such as bar graphs, line charts, and pie charts. Interactive dashboards are then developed by combining multiple visualizations, allowing users to explore data, apply filters, and gain insights easily.

The data flows from the data source layer through the ETL process, where it is cleaned and stored. OLAP operations are then applied for analysis, and finally, Apache Superset presents the results in the form of interactive dashboards. This architecture ensures efficient data processing, analysis, and visualization for better decision-making.

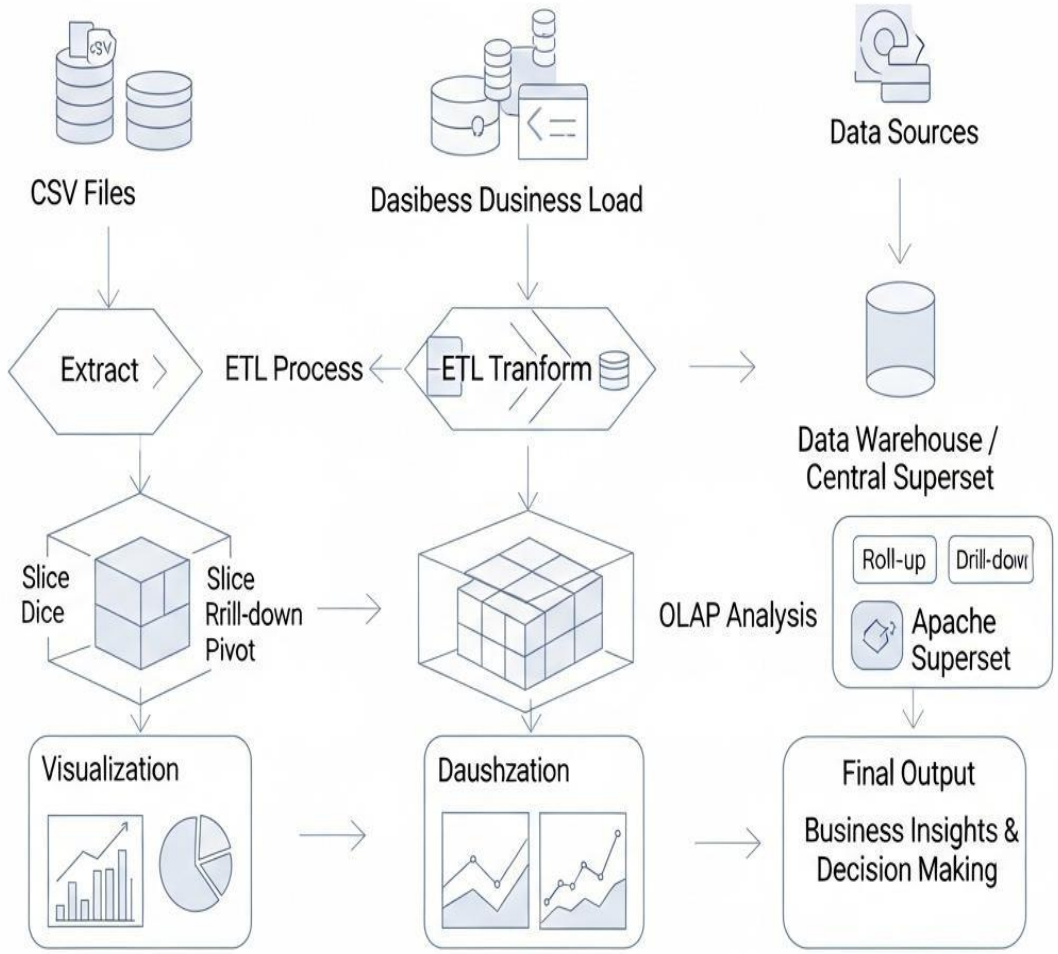

Figure 1: System Architecture — Data Flow from Source to Visualization Layer

The above diagram illustrates the three-layer architecture of the system, depicting the flow of data from raw sources through the ETL and OLAP processing layer to the final visualization layer in Apache Superset.

8. IMPLEMENTATION

The project "Business Data Analysis using OLAP and Apache Superset" is implemented using a structured approach that integrates data processing, analytical techniques, and visualization tools.

Initially, the dataset — comprising sales, product, and regional data — is collected from various sources in formats such as CSV files or databases. Since the data may be inconsistent and unstructured, it is processed using the ETL (Extract, Transform, Load) methodology.

In the Extract phase, data is gathered from multiple sources. In the Transform phase, the data is cleaned, filtered, and converted into a consistent format suitable for analysis. In the Load phase, the processed data is stored in a centralized database such as MySQL or PostgreSQL.

After storing the data, OLAP (Online Analytical Processing) techniques are applied. The data is organized into multidimensional structures known as cubes, where dimensions (such as time, location, and product) and measures (such as sales and revenue) are defined.

Various OLAP operations are implemented to analyze the data:

- Slice — to focus on a specific dimension
- Dice — to filter multiple dimensions
- Roll-up — to summarize data at a higher level
- Drill-down — to view detailed data
- Pivot — to change the data perspective

For visualization, Apache Superset is connected to the database. Datasets are created within Superset, and different types of charts such as bar charts, line charts, and pie charts are generated. These visualizations are combined into interactive dashboards, allowing users to explore data, apply filters, and gain meaningful insights.

Overall, the implementation demonstrates how raw data can be transformed into valuable insights using ETL processes, OLAP techniques, and modern business intelligence tools like Apache Superset.

9. RESULTS AND OUTPUT

The project successfully demonstrates the effective implementation of OLAP techniques and Apache Superset for business data analysis and visualization.

One of the key outcomes of the project is the successful integration of the ETL (Extract, Transform, Load) process, which enabled the collection, cleaning, and consolidation of data from multiple sources into a centralized database. This ensured that the data used for analysis was consistent, accurate, and reliable.

The application of OLAP operations such as slice, dice, roll-up, drill-down, and pivot allowed efficient multidimensional analysis of the dataset. These operations helped in exploring data from different perspectives, identifying trends, and performing both detailed and summarized analyses.

Using Apache Superset, interactive dashboards and visualizations were created, including bar charts, line graphs, and pie charts. These dashboards made it easier to understand complex data and provided clear insights into business performance, such as sales trends, product comparisons, and regional analysis.

The system was tested with different datasets and scenarios, and it performed efficiently in handling large volumes of data and generating visual outputs without significant delay.

9.1 Dataset View

The dataset used in this project consists of structured business data including product details, sales figures, time periods, and regional information. Each record represents a transaction with attributes such as product name, date, location, and revenue. The data is processed using the ETL method to ensure consistency and accuracy. It is then used for OLAP analysis and visualized using Apache Superset to generate meaningful insights.

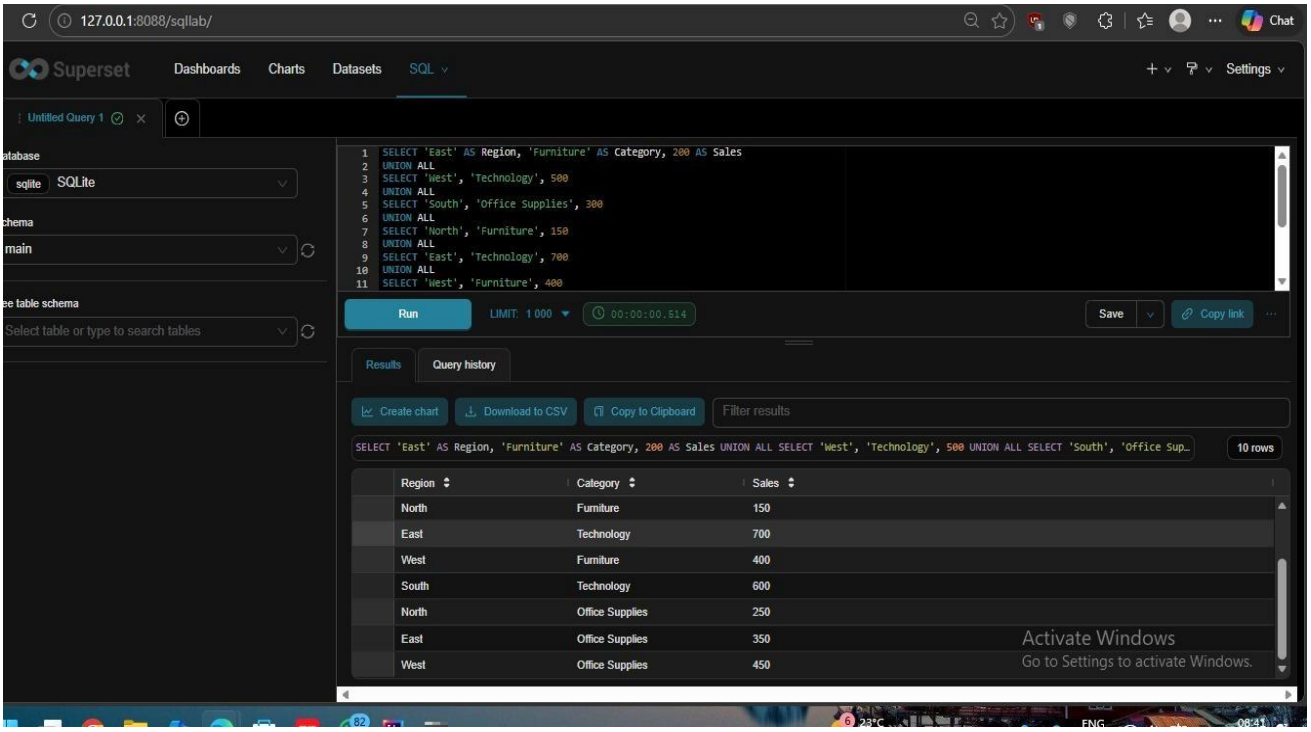

Figure 2: Dataset View — Structured Business Data Loaded in Apache Superset

The figure above displays the tabular view of the business dataset loaded into the system, showing key attributes such as product category, region, date, and sales values used for OLAP analysis.

9.2 Roll-Up Operation

Roll-up is a fundamental OLAP operation used to summarize or aggregate data by moving up in a dimensional hierarchy. It helps in converting detailed data into a higher-level view, making it easier to analyze overall trends and patterns. In a data cube, roll-up reduces the level of detail by grouping data based on higher-level dimensions.

For example, sales data collected at the city level can be aggregated to the state or country level, or data recorded daily can be summarized into monthly or yearly values. This operation is useful for decision-making as it provides a concise overview of large datasets. Roll-up is the inverse of the drill-down operation, which focuses on providing more granular data.

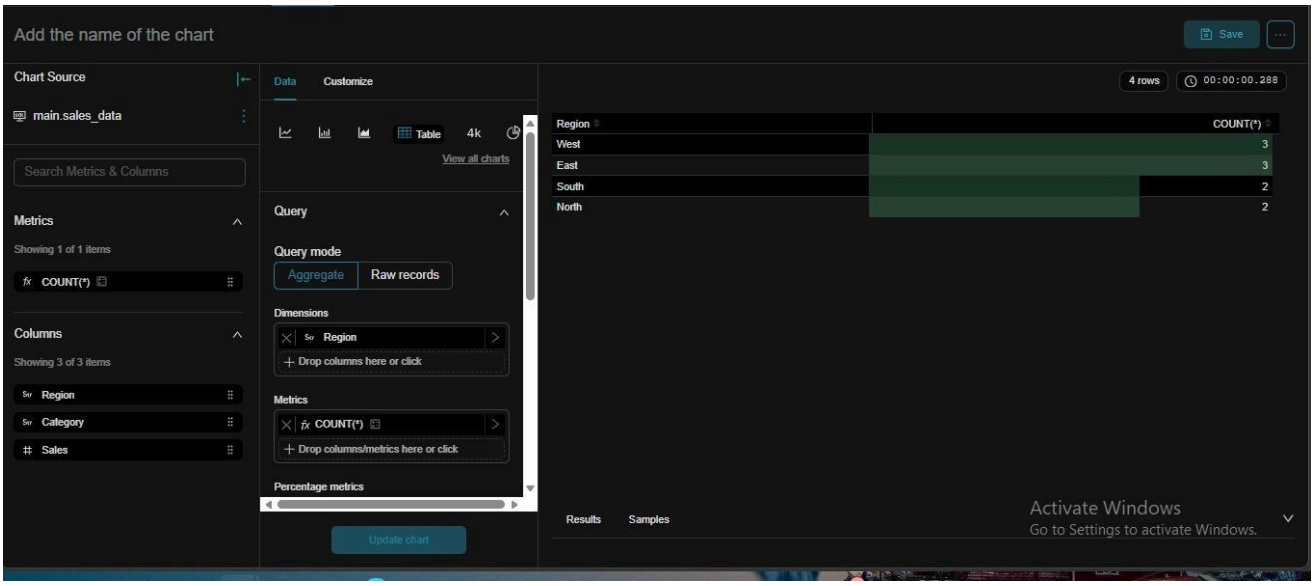

Figure 3: Roll-Up Operation — Aggregating Data by Higher-Level Dimensions

The figure above illustrates the roll-up operation applied to the dataset, where granular transaction-level data is aggregated to a higher-level summary, enabling quick identification of overall business trends.

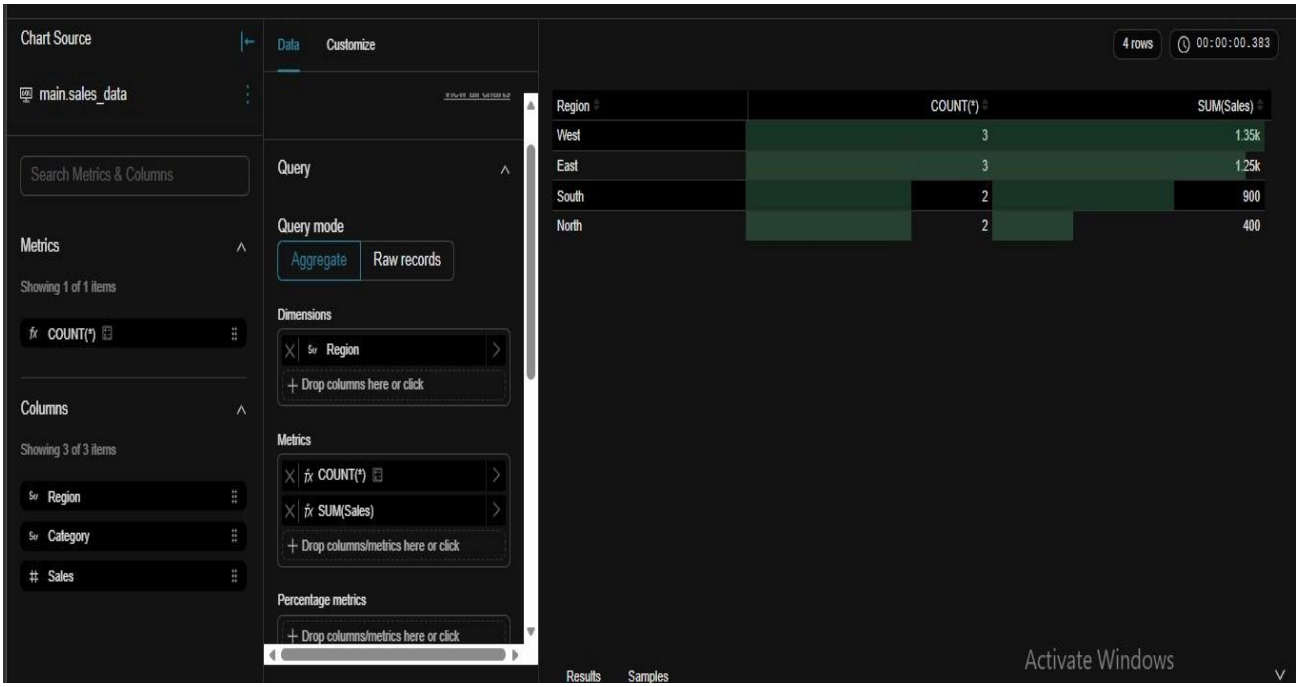

Figure 4: Roll-Up Result — Summarized Sales Data across Regions

This output shows the result of the roll-up operation, presenting aggregated sales figures grouped by region and time period to provide a high-level overview of business performance.

9.3 Drill-Down Operation

Drill-down is an OLAP operation used to view data in greater detail by moving down the dimensional hierarchy. It helps users analyze data at a lower level of granularity by breaking down summarized data into finer details. For example, sales data shown at the yearly level can be drilled down to view quarterly, monthly, or even daily data. Similarly, data at the country level can be drilled down to the state or city level, enabling deeper investigation into specific segments.

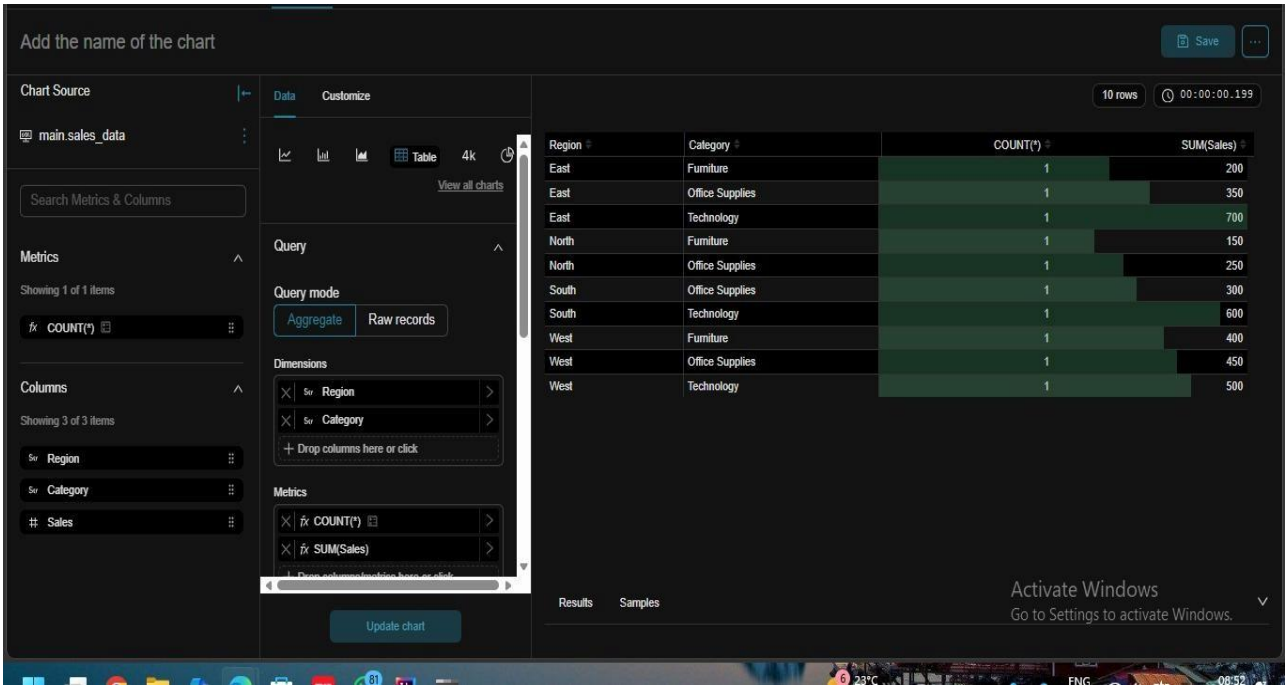

Figure 5: Drill-Down Operation — Viewing Detailed Data at Lower Hierarchy Levels

The figure above demonstrates the drill-down operation, where summarized annual sales data is broken down into quarterly and monthly details, allowing analysts to investigate performance variations at a finer level of granularity.

9.4 Slice Operation

Slice is an OLAP operation used to extract a specific subset of data from a multidimensional data cube by selecting a single dimension. It helps in focusing on a particular value of one dimension while analyzing the remaining data. For example, selecting data for a specific year (e.g., 2024) from a dataset allows analysis of all regions and products for that year only. This reduces the complexity of the data and makes it easier to study a particular aspect of the dataset in detail.

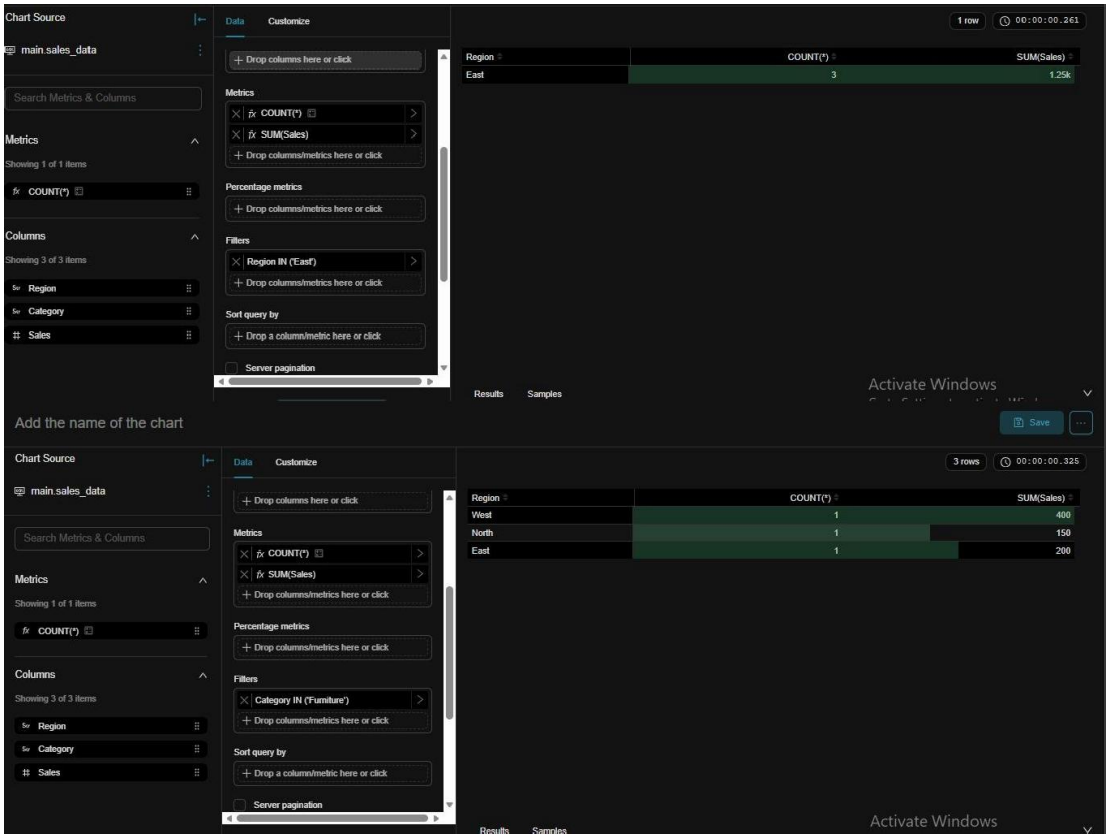

Figure 6: Slice Operation — Extracting a Subset Based on a Single Dimension

The figure above depicts the slice operation, where the dataset is filtered along a single dimension — in this case, a specific time period — to isolate and analyze relevant data without the complexity of the full dataset.

9.5 Dice Operation

Dice is an OLAP operation used to extract a specific subset of data from a multidimensional data cube by applying filters on two or more dimensions simultaneously. It allows users to analyze data by selecting multiple values across different dimensions. For example, selecting data for regions such as East and West, categories such as Furniture and Technology, and a specific time period creates a smaller sub-cube focused on these conditions. This operation enables more detailed and targeted analysis compared to the slice operation.

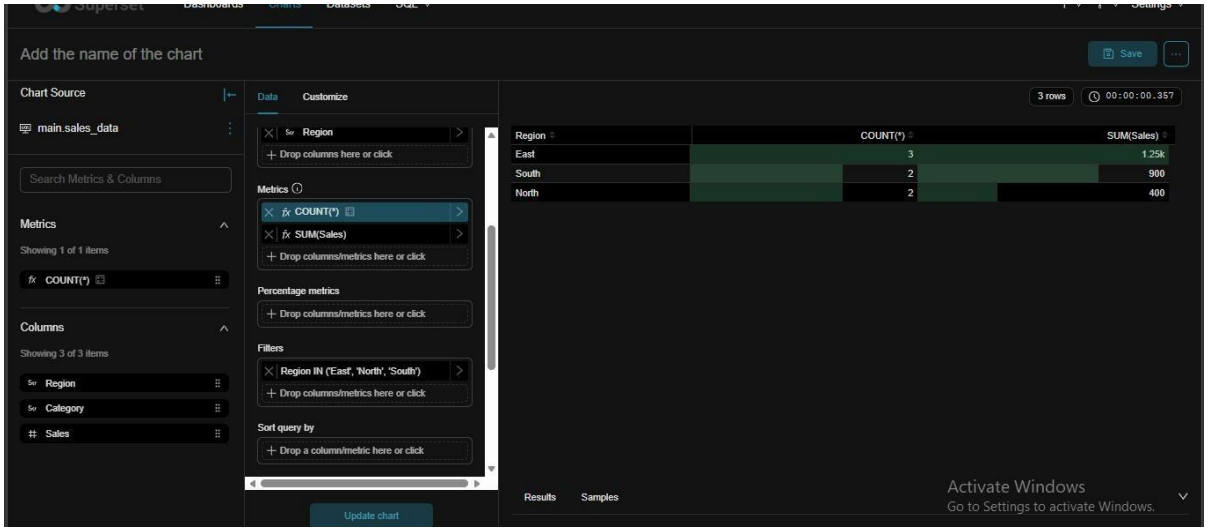

Figure 7: Dice Operation — Filtering Data Across Multiple Dimensions

The figure above illustrates the dice operation, where data is filtered across multiple dimensions — including region, product category, and time period — to produce a focused sub-cube suitable for detailed targeted analysis.

9.6 Bar Graph Visualization

A bar graph is a type of data visualization used to represent data using rectangular bars, where the length or height of each bar corresponds to the value it represents. It is commonly used to compare different categories of data in a clear and straightforward manner. In this project, bar graphs are generated in Apache Superset to compare values such as sales figures across different regions or product categories. Each bar represents a category, and its size indicates the magnitude of the corresponding data value.

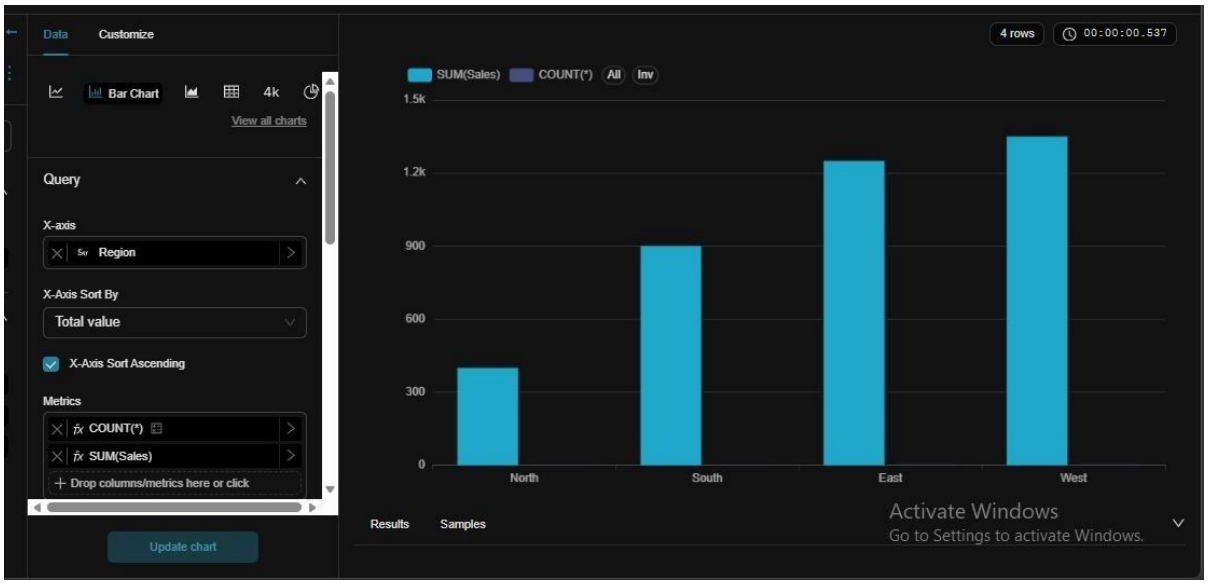

Figure 8: Bar Graph — Sales Comparison Across Regions and Product Categories

The bar graph above visualizes sales data across various product categories and regions, enabling quick comparison of performance metrics and identification of top-performing segments within the business dataset.

9.7 Pie Chart Visualization

A pie chart is a circular data visualization that represents the proportional contribution of each category to a whole. The entire circle represents the total dataset, and each slice corresponds to a category, with the size of the slice indicating the relative proportion of that category. In this project, pie charts are generated in Apache Superset to display the distribution of sales across different regions or product categories, providing an intuitive view of how the total is composed.

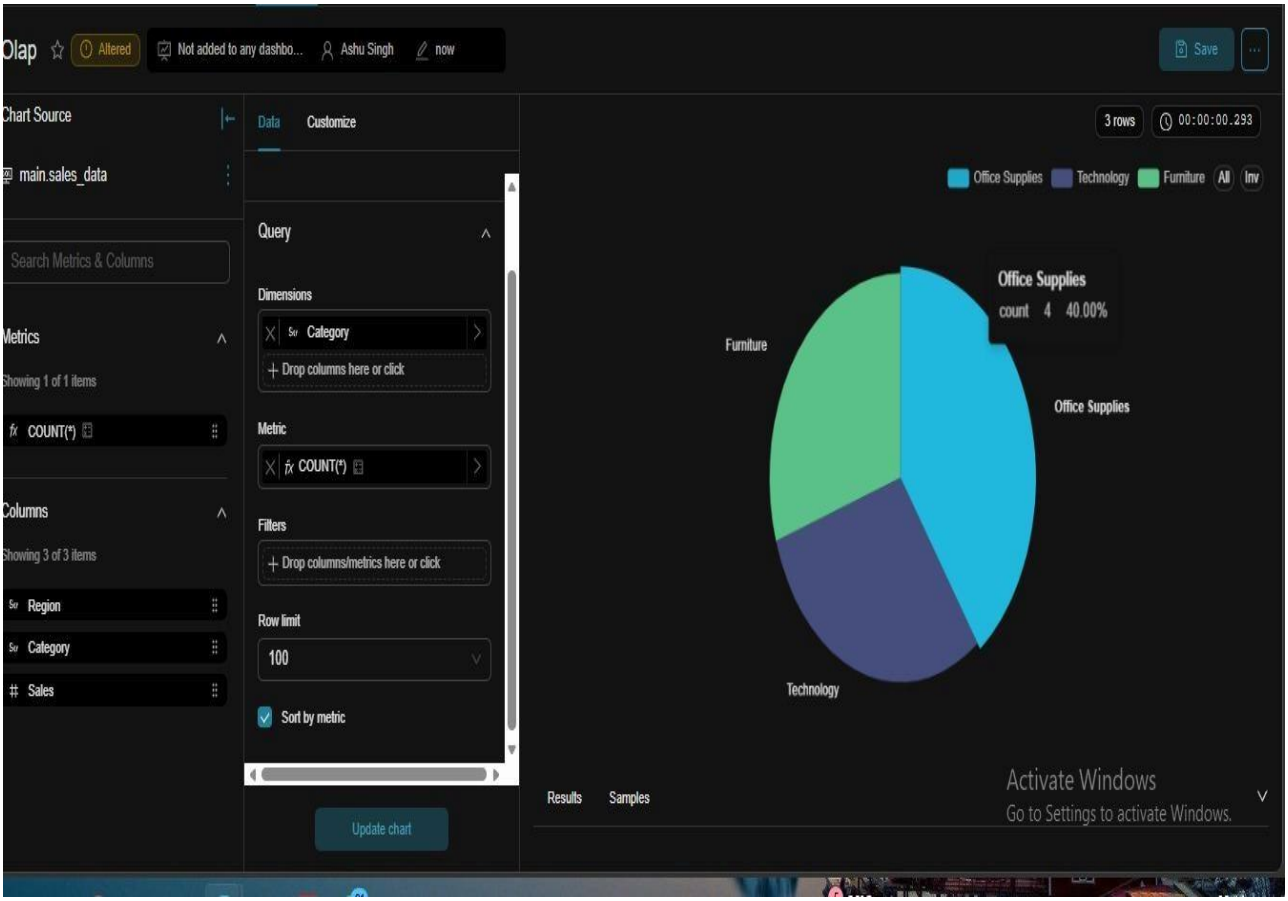

Figure 9: Pie Chart — Distribution of Sales by Product Category

The pie chart above illustrates the proportional share of each product category in total sales, helping stakeholders quickly identify dominant and underperforming categories within the dataset.

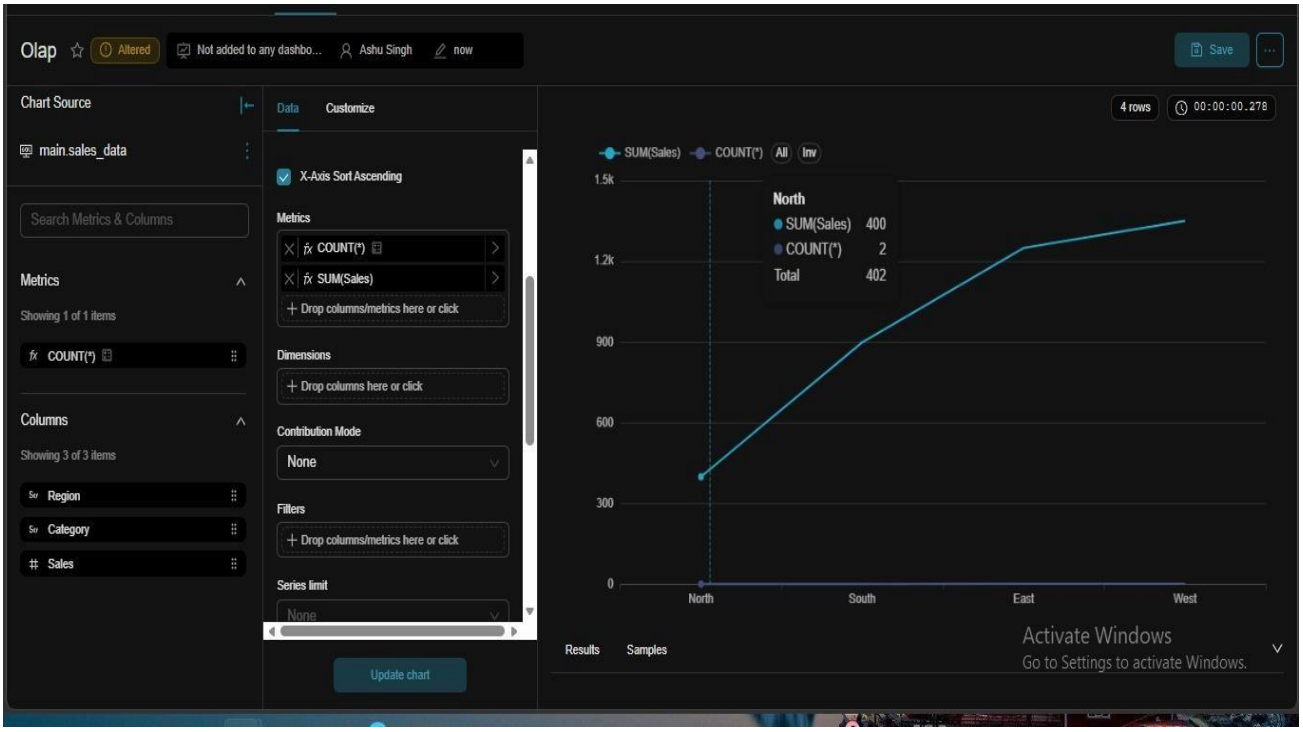

Figure 10: Pie Chart — Regional Sales Distribution

This second pie chart presents the distribution of sales across geographic regions, providing a clear picture of regional contributions and supporting location-based business decision-making.

9.8 Pivot Operation

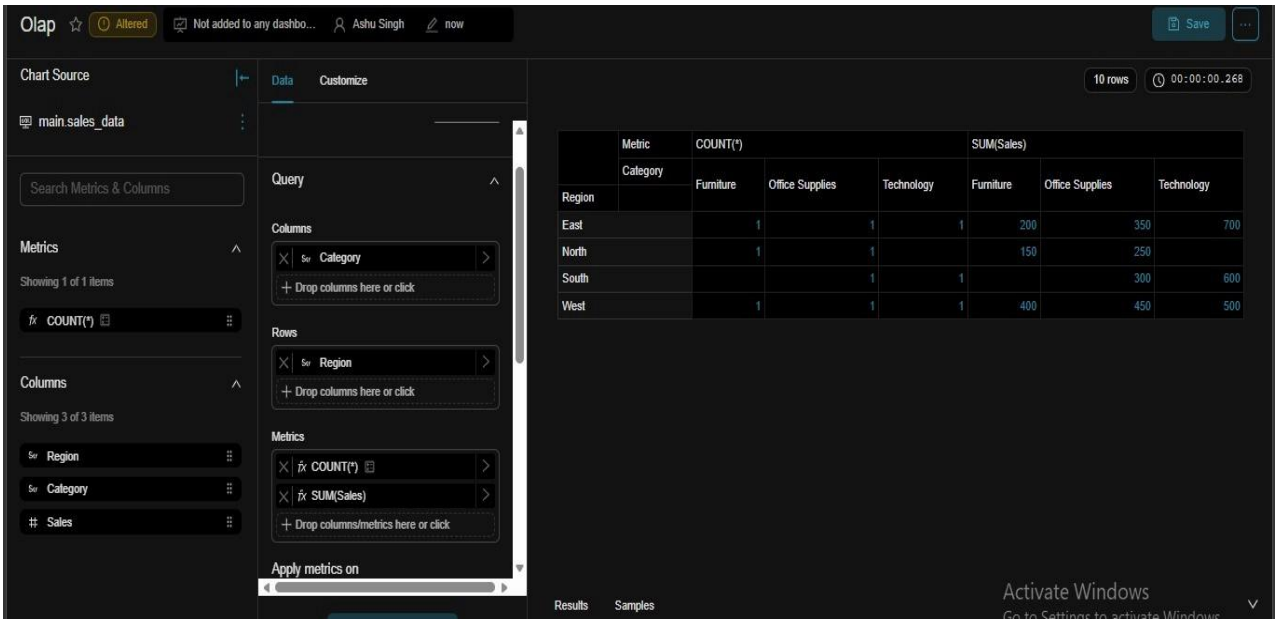

	Metric	COUNT(*)			SUM(Sales)		
	Category	Furniture	Office Supplies	Technology	Furniture	Office Supplies	Technology
Region							
East		1	1	1	200	350	700
North		1	1		150	250	
South			1	1		300	600
West		1	1	1	400	450	500

Figure 11: Pivot Operation — Rotating Data Perspective for Multi-Angle Analysis

The figure above demonstrates the pivot operation, where rows and columns of the data cube are rearranged to present the same dataset from a different analytical perspective, facilitating better comparison and deeper understanding.

Pivot, also known as rotation, is an OLAP operation used to change the orientation or perspective of data in a multidimensional data cube. It allows users to rearrange rows and columns to view data from different angles without altering the underlying data. For example, sales data can be viewed with regions as rows and products as columns, or vice versa. This operation helps in better comparison and understanding of data by providing multiple viewpoints for analysis.

10. CONCLUSION

The project "Business Data Analysis using OLAP and Apache Superset" was successfully developed and implemented. It demonstrates how modern data analysis techniques and business intelligence tools can be combined to transform raw data into meaningful insights.

The system effectively utilizes the ETL (Extract, Transform, Load) process to collect, clean, and organize data from multiple sources. By applying OLAP operations such as slice, dice, roll-up, drill-down, and pivot, the project enables efficient multidimensional analysis of data.

The integration of Apache Superset for data visualization enhances the user experience by presenting data in the form of interactive charts and dashboards. This makes it easier to understand trends, patterns, and key performance indicators.

The project also provided practical knowledge of data processing, analytical techniques, and visualization tools. It highlights the importance of data-driven decision-making in modern organizations.

Overall, this project successfully achieves its objectives and demonstrates how OLAP and business intelligence tools can be used to improve data analysis, support better decision-making, and enhance overall business performance.

11. FUTURE SCOPE

The following enhancements and extensions are proposed for future development of this system:

- Integrate real-time data streaming to enable live data analysis and dynamic dashboards.
- Implement role-based access control to ensure data security and user-specific dashboard views.
- Enhance dashboard features with more advanced visualizations and customization options.
- Integrate the system with cloud platforms such as AWS or Azure for improved scalability and performance.
- Enable automated reporting and alert systems based on key performance indicators (KPIs).

12. REFERENCES

[1] A. Silberschatz, H. F. Korth, and S. Sudarshan, Database System Concepts, 7th ed. New York, NY, USA: McGraw-Hill Education, 2019.

[2] R. Elmasri and S. B. Navathe, Fundamentals of Database Systems, 7th ed. Pearson Education, 2016.

[3] S. Chaudhuri and U. Dayal, "An Overview of Data Warehousing and OLAP Technology," ACM SIGMOD Record, vol. 26, no. 1, pp. 65–74, 1997.

[4] W. H. Inmon, Building the Data Warehouse, 4th ed. Wiley, 2005.

[5] R. Kimball and M. Ross, The Data Warehouse Toolkit: The Definitive Guide to Dimensional Modeling, 3rd ed. Wiley, 2013.

[6] Apache Software Foundation, "Apache Superset Documentation." [Online]. Available: https://superset.apache.org

[7] Microsoft, "OLAP (Online Analytical Processing)," Microsoft Docs. [Online]. Available: https://learn.microsoft.com

[8] Oracle Corporation, "Introduction to OLAP," Oracle Documentation. [Online]. Available: https://docs.oracle.com